

Exploring concurrent and stateless evolutionary algorithms

Seeking the best communication strategies in a communicating sequential process framework

Anonymous Author 1

Anonymous institute 1
Dublin, Ohio
aa1@ai1.com

Anonymous Author 3

Anonymous institute 2
Dublin, Ohio
2aa@ai2.com

Anonymous Author 2

Anonymous institute 2
Dublin, Ohio
2aa@ai2.com

Anonymous Author 4

Anonymous institute 2
Dublin, Ohio
2aa@ai2.com

ABSTRACT

Concurrent algorithms use channels for communication, which implies that communication is an integral part of them, so some attention must be devoted to its design. In the design of concurrent evolutionary algorithms, there are several options that can be used for performing this communication. In this paper we will explore how communication overhead can be reduced, and how it influences scaling. The evolutionary algorithm will use a concurrent language, and leverage its capabilities. Eventually, we will try to prove how concurrent version of algorithms offer a good option to leverage the multi-threaded and multi-core capabilities of modern computers.

CCS CONCEPTS

• **Theory of computation** → **Concurrent algorithms**; • **Computing methodologies** → **Genetic algorithms**; • **General and reference** → *Performance*;

KEYWORDS

Concurrency, Concurrent evolutionary algorithms, performance evaluation, algorithm design

ACM Reference Format:

Anonymous Author 1, Anonymous Author 2, Anonymous Author 3, and Anonymous Author 4. 2019. Exploring concurrent and stateless evolutionary algorithms: Seeking the best communication strategies in a communicating sequential process framework. In *Proceedings of the Genetic and Evolutionary Computation Conference 2019 (GECCO '19)*. ACM, New York, NY, USA, 6 pages. <https://doi.org/10.1145/nnnnnnn.nnnnnnn>

1 INTRODUCTION

Despite the emphasis on hardware-based techniques such as cloud computing or GPGPU, there are not many papers [15] dealing with creating concurrent evolutionary algorithms that work in a single computing node or that extend seamlessly from single to many computers.

The concurrent programming paradigm (or concurrency oriented programming [2]) is characterized by the presence of programming constructs for managing processes like first-class objects. That is, with native operators for acting upon them and the possibility of using them as parameters or as return values of a function. The latter yields to changes in the implementation of concurrent algorithms due to the direct mapping between patterns of communications and processes with language expressions; on one hand it becomes simpler since the language provides an abstraction for communication, on the other hand it changes the paradigm for implementing algorithms, since these new communication constructs have to be taken into account.

Moreover, concurrent programming adds a layer of abstraction over the parallel facilities of processors and operating systems, offering a high-level interface that allows the user to program modules of code to be executed in parallel threads [1].

Different languages offer different concurrency strategies depending on how they deal with shared state, that is, data structures that could be accessed from several processes. In this regard, there are two major fields (with some other variations):

- Actor-based concurrency [17] totally eliminates shared state by introducing a series of data structures called *actors* that store state and can mutate it locally.
- Process calculi or process algebra is a framework to describe systems that work with independent processes that interact between them using channels. One of the best known is called the *communicating sequential processes* (CSP) methodology [9], which is effectively stateless, with different processes reacting to a channel input without changing state, and writing to these channels. Unlike actor based concurrency, which keeps state local, in this case per-process state is totally eliminated, with all computation state managed as messages in a channel.
- Other, less well known models using, for instance, tuple spaces [7].

Most modern languages, however, follow the CSP abstraction, and it has become popular since it fits well other programming paradigms, like reactive and functional programming, and allows for a more efficient implementation, with less overhead, and with well-defined primitives. This is why we will use it in this paper for creating new evolutionary algorithms that live *natively* in these environments and can thus be implemented easily in this kind of

Permission to make digital or hard copies of part or all of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for third-party components of this work must be honored. For all other uses, contact the owner/author(s).

GECCO '19, July 13–17, 2019, Prague, Czech Republic

© 2019 Copyright held by the owner/author(s).

ACM ISBN 978-x-xxxx-xxxx-x/YY/MM...\$15.00

<https://doi.org/10.1145/nnnnnnn.nnnnnnn>

languages. We have chosen Perl 6, although other languages such as Go, are feasible alternatives.

In previous papers [22, 23] we designed an evolutionary algorithm that fits well this architecture and explored its possibilities. That initial exploration showed that a critical factor within this algorithmic model is communication between threads; therefore designing efficient messages is high-priority to obtain a good algorithmic performance and scaling. In this paper, we will test several communication strategies: a lossless one that compresses the population, and a lossy one that sends a representation of population gene-wise statistics.

2 STATE OF THE ART

The problem of communication/synchronization between processes, which nowadays is accomplished by means of threads, has been the subject of long-standing study. One of the best efforts to formalize and simplify that matter is Hoare's *Communicating Sequential Processes* [9], that proposes an interaction description language which is the theoretical support for many libraries and modern programming languages. In such model, concurrent programs communicate on the basis of *channels*, a sort of pipe buffer used to interchange messages between the different processes or threads, either asynchronously or synchronously. Languages such as Go and Perl 6 implement this concurrency model (the latter including additional mechanisms such as *promises* or low-level access to the creation of threads).

The fact that messages have to be processed without secondary effects and that actors do not share any kind of state makes concurrent programming specially fit for functional languages or languages with functional features; this has made this paradigm specially popular for late cloud computing implementations; however, its presence in the EA world is not so widespread, although some efforts have lately revived the interest for this kind of paradigm [20]. Several years ago it was used in Genetic Programming [4, 10, 26] and recently in neuroevolution [18] but its occurrence in EA, despite being scarce in the previous years [8], has experimented a certain rise lately with papers such as [25] which perform program synthesis using the functional programming features of the Erlang language [3] for building an evolutionary multi-agent system.

Regarding functional languages, Erlang and Scala have embraced the actor model of concurrency and get excellent results in many application domains; Clojure is another one with concurrent features such as promises/futures, Software Transaction Memory and agents; Kotlin [19] has been recently used for implementing a functional evolutionary algorithm framework.

On the other hand, Perl 6 [?] uses different concurrency models, varying from implicit concurrency using a particular function that automatically parallelizes operations on iterable data structures, to explicit concurrency using threads. These both types of concurrency will be considered in this study whilst using the Algorithm: :Evolutionary: :Simple, a Perl 6 library introduced in last year's version of this very same conference [16].

Earlier efforts to study the issues of concurrency in EA are worth mentioning. For instance, the EvAg model [11] resorts to the underlying platform scheduler to manage the different threads of execution of the evolving agents; in this way the model scaled-up

seamlessly to take full advantage of CPU cores. In the same avenue of measuring scalability, experiments were conducted in [13] comparing single and a dual-core processor concurrency achieving near linear speed-ups. The latter was later on extended in [14] by scaling up the experiment to up to 188 parallel machines, reporting speed-ups up to 960×, nearly four times the expected linear growth in the number of machines (when local concurrency were not taken into account). Other authors have addressed explicitly multi-core architectures, such as Tagawa [21] which used shared memory and a clever mechanism to avoid deadlock. Similarly, [12] used a message-based architecture developed in Erlang, separating GA populations as different processes, although all communication taking place with a common central thread.

In previous papers [22, 24], we presented a proof of concept of the implementation of stateless evolutionary algorithms using Perl 6, based on a single channel model communicating threads for population evolving and mixing. In addition, we studied the effect of running parameters such as *generation gap* (similar to the concept of *time to migration* in parallel evolutionary algorithms) and population size, realizing that the choice of parameters may have a strong influence at the algorithmic level, but also at the implementation level, in fact affecting the actual wallclock performance of the EA.

3 EXPERIMENTAL SETUP




Figure 1: General scheme of operation of channels and thread groups.

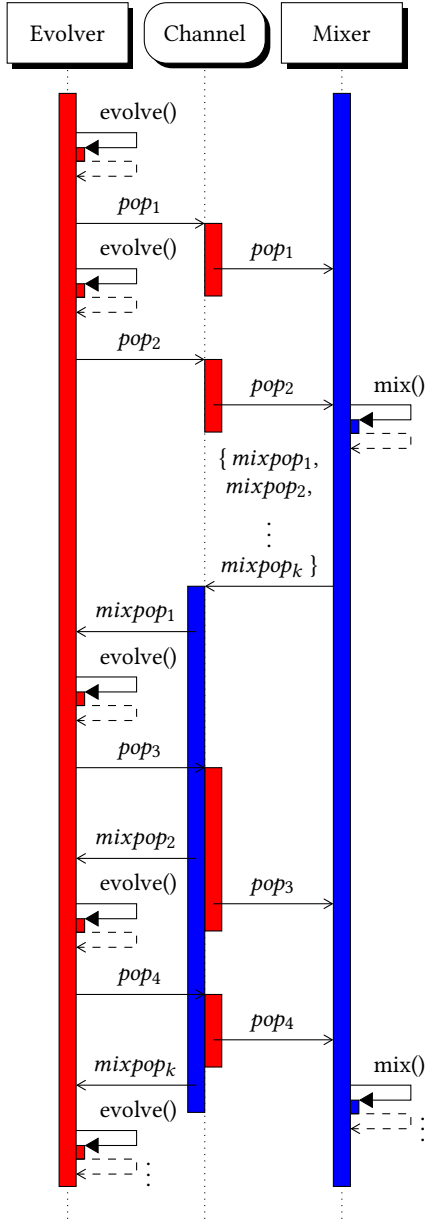


Figure 2: Schematic of communication between threads and channels for concurrent EAs. The two central bars represent the channel, and color corresponds to the *main* function they perform; blue for mixer, red for evolver. As the figure shows, the evolver threads always read from the mixer channel, and always write to the evolver channel.

The baseline we are coming from is similar to the one used in previous experiments [22]. Our intention was to create a system that was not functionally equivalent to a sequential evolutionary algorithms, and that followed the principle of communicating sequential processes. In this kind of methodology, we will have processes (or threads) communicating state through channels. Every

process itself will be stateless, reacting to the presence of messages in the channels it is listening to and sending result back to them, without changing state.

As in the previous papers, [23], we will use two groups of threads and two channels. We will see them in turns.

The two groups of threads perform the following functions:

- The *evolutionary* threads will be the ones that will be in principle running the evolutionary algorithm.
- The *mixing* threads will *mix* populations, and create new ones as a mixture of them.

The two channels carry messages that are equivalent to populations, but they do so in a different way:

- The *evolutionary* channel will be used for carrying non-evolved, or generated, populations.
- The *mixer* channel will carry, *in pairs*, evolved populations.

These will be connected as shown in Figure 1. The evolutionary group of threads will read only from the evolutionary channel, evolve for a number of generations, and place result in the mixer channel; the mixer group of threads will read only from the mixer channel, in pairs. From every pair, a random element is put back into the mixer channel, and a new population is generated and sent back to the evolutionary channel. The main objective of using two channels is to avoid deadlocks; the fact that one population is written always back to the mixer channel avoids starvation in the channel. How this runs in practice is shown in Figure 2, where the timeline of the interchange of messages between the evolver and mixer threads and evolver and mixer channels is clarified.

One of the problems of the baseline configuration was that communication took a great amount of time, adding some overhead to the algorithm. The *message* consisted of the whole population, and the size increased with population size, obviously. This tipped the balance between communication and computation towards communication, so that the more threads, the more communication was taking place. Our first intention in this paper was to slim down messages so that they took less bandwidth (or memory) and less time to send and process. In general, this strategy also can be framed in the context of migration strategies, since that is the most similar thing in the context of parallel algorithms. In parallel algorithms, an adequate selection of migration strategies, balancing exploration and exploitation, is the key to achieving high performance, as indicated in [5]. In the context of concurrent evolutionary algorithms we will talk about *population messages*, but their effect is going to be similar. For this paper, we introduced two different messaging strategies:

- One we have called *EDA*, or estimation of distribution algorithm, whose basic idea is that the population message will contain the probability distribution over each gene. In this sense, this strategy is similar to the one presented by de la Ossa et al. in [6]. Not being an estimation of distribution algorithm *per se*, since the evolutionary thread runs a canonical genetic algorithm, when the message is being composed, every (binary) gene of the 25% best individuals in the population is examined, and an array with the probabilities for each gene is sent to the mixer thread. The *mixer* thread, in turn, just takes randomly one probability from each of the two *populations* (actually, distributions), instead of working

on individuals. While in the baseline strategy the selection took place in the mixer thread, that eliminated half the population, in this case the selection takes place when composing the message, since just the 25% best individuals are selected to compute the probability distribution of genes. When the evolver thread reads the message, it rebuilds the population based on this distribution.

- The second is called *compress*, and it simply bit-packs the population, without the fitness, into a message which uses 1 bit per individual, and then 64 bits, or simply 8 bytes, to transmit a single individual in the population. This strategy is equivalent to the baseline, except it introduces an additional step of evaluating the population when mixing and receiving it from the evolutionary channel. It is hoped that this additional evaluation overhead does compensate the communication overhead that is eliminated.

In the same way we did in our previous papers, first we will have to evaluate these new strategies compared with baseline; since the overhead will be different depending on the computation time, which in this case is regulated by the number of generations that are going to be performed by every thread, we will first perform an experiment changing that. We will call this parameter the *generation gap*, implying that's the gap between receiving a message and activating the thread and sending it, deactivating it. Besides, what we want to find out in these set of experiments is what is the generation gap that gives the best performance in terms of raw time to find a solution, as well as the best number of evaluations per second.

Additionally, it is impossible to know from first principles if this setup is the only possible. We have to heuristically explore other possibilities; in this case, we will explore another messaging strategy called *no writeback*, or *nw*, where the mixer thread, instead of sending one of the individuals back again to the mixer thread, sends it to the evolver channel, where it will undergo an additional round of evolution. The main difference between these two strategies is twofold:

- The mixer channel can be empty for some time, since it is not always holding at least one population message (written back every time it is activated). This might lead to *starvation* of the mixer thread, but in fact it will not take long since it is going to be processed immediately by the evolver thread.
- Every population is mixed just once, which might lead to improvements in the algorithm.

4 EXPERIMENTAL RESULTS

The experiments have been prepared by using OneMax function with 64 bits. This function was chosen primarily since it was the one used in previous experiments, but also because it is a classical benchmark, it can be easily programmed in Perl 6, and allowed us to focus in what we are interested in, the design of the concurrent evolutionary algorithm itself. We used the open source `Algorithm::Evolutionary::Simple` Perl 6 module for the evolutionary part, and wrote the different scripts in the same language. The latest version, compiled from source, of Perl 6 was used, and experiments were performed in a Intel(R) Core(TM) i7-4770 CPU at 3.40GHz running Ubuntu 14.04 server. All scripts have a free

Table 1: Parameters used to explore the generation gap

Parameter	Value
Evolver threads	2
Mixer threads	1
Generations	8,16,32
Population size	256
Initial populations	3
Bits	64

license and have been released in GitHub, where the data generated by every single experiment is also hosted.

In these experiments, the parameters used are shown in Table 1. Two evolver threads are needed, at least, to avoid starvation of the mixer thread. The generation gap was checked for those values in all cases, although in some cases we extended it to 4 and 64 generations. The population was sized in previous papers using the bisection method, and the number of initial populations created and sent to the evolver channel was also designed to avoid starvation; that way, as soon as the first two populations are evaluated simultaneously by the two threads and sent to the mixer channel, this channel will always hold a population that will combine with a fresh one coming from either of the mixer threads.

5 CONCLUSIONS

In this paper we have set out to explore the interaction between the generation gap and the algorithmic parameters in a concurrent and stateless evolutionary algorithm. From the point of view of the algorithm, increasing the generation gap favors exploitation over exploration, which might be a plus in some problems, but also decreases diversity, which might lead to premature convergence; in a parallel setting, this will make the algorithm need more evaluations to find a solution. The effect in a concurrent program goes in the opposite direction: by decreasing communication, the amount of code that can be executed concurrently increases, increasing performance. Since the two effects cancel out, in this paper we have used an experimental methodology to find out what is the combination that is able to minimize wallclock time, which is eventually what we are interested in by maximizing the number of evaluations per second while, at the same time, increasing by a small quantity the number of evaluations needed to find the solution.

For the specific problem we have used in this short paper, a 64-bit onemax, the generation gap that is in that area is 16. The time to communication for that specific generation gap is around 2 seconds, since 16 generations imply 4096 evaluations and evaluation speed is approximately 2K/s. This gives us a ballpark of the amount of computation that is needed for concurrency to be efficient. In this case, we are sending the whole population to the communication channel, and this implies a certain overhead in reading, transmitting and writing. Increasing the population size also increases that overhead.

We can thus deduce that the amount of computation, for this particular machine, should be on the order of 2 seconds, so that it effectively overcomes the amount of communication needed. This amount could be played out in different way, for instance by




Figure 3: Screenshot using the htop utility of the used machine running two experiments at the same time. As it can be seen, all processors are kept busy, with a very high load average.

increasing the population; if the evaluation function takes more time, different combinations should be tested so that no message is sent unless that particular amount of time is reached.

With these conclusions in mind, we can set out to work with other parameters, such as population size or number of initial populations, so that the loss of diversity for bigger population sizes is overcome. Also we have to somehow overcome the problem of the message size by using a statistical distribution of the population, or simply other different setup. This is left as future work.

ACKNOWLEDGMENTS

Acknowledgements taking this much space

REFERENCES

- [1] Gregory R Andrews. 1991. *Concurrent programming: principles and practice*. Benjamin/Cummings Publishing Company San Francisco.
- [2] Joe Armstrong. 2003. Concurrency Oriented Programming in Erlang. (2003). <http://ll2.ai.mit.edu/talks/armstrong.pdf>
- [3] Adam D Barwell, Christopher Brown, Kevin Hammond, Wojciech Turek, and Aleksander Byrski. 2017. Using program shaping and algorithmic skeletons to parallelise an evolutionary multi-agent system in Erlang. *Computing and Informatics* 35, 4 (2017), 792–818.
- [4] Forrest Briggs and Melissa O'Neill. 2008. Functional genetic programming and exhaustive program search with combinator expressions. *Int. J. Know.-Based Intell. Eng. Syst.* 12, 1 (Jan. 2008), 47–68. <http://dl.acm.org/citation.cfm?id=1375341>.
- [5] Erick Cantú-Paz. 1999. Migration Policies and Takeover Times in Genetic Algorithms. In *Proceedings of the 1st Annual Conference on Genetic and Evolutionary Computation - Volume 1 (GECCO '99)*. Morgan Kaufmann Publishers Inc., San Francisco, CA, USA, 775–775. <http://dl.acm.org/citation.cfm?id=2933923.2934003>
- [6] Luis delaOssa, José A. Gámez, and José M. Puerta. 2004. Migration of Probability Models Instead of Individuals: An Alternative When Applying the Island Model to EDAs. In *Parallel Problem Solving from Nature - PPSN VIII*, Xin Yao, Edmund K. Burke, José A. Lozano, Jim Smith, Juan Julián Merelo-Guervós, John A. Bullinaria, Jonathan E. Rowe, Peter Tiño, Ata Kabán, and Hans-Paul Schwefel (Eds.). Springer Berlin Heidelberg, Berlin, Heidelberg, 242–252.
- [7] David Gelernter. 1985. Generative communication in Linda. *ACM Transactions on Programming Languages and Systems (TOPLAS)* 7, 1 (1985), 80–112.
- [8] John Hawkins and Ali Abdallah. 2001. A Generic Functional Genetic Algorithm. In *Proceedings of the ACS/IEEE International Conference on Computer Systems and Applications (AICCSA '01)*. IEEE Computer Society, Washington, DC, USA, 11–. <http://dl.acm.org/citation.cfm?id=872017.872197>
- [9] C. A. R. Hoare. 1978. Communicating sequential processes. *Commun. ACM* 21, 8 (Aug. 1978), 666–677. <https://doi.org/10.1145/359576.359585>
- [10] Lorenz Huelsbergen. 1996. Toward simulated evolution of machine-language iteration. In *Proceedings of the First Annual Conference on Genetic Programming (GECCO '96)*. MIT Press, Cambridge, MA, USA, 315–320. <http://dl.acm.org/citation.cfm?id=1595536.1595579>
- [11] Juan-Luis Jiménez-Laredo, A. E. Eiben, Maarten van Steen, and Juan-Julian Merelo-Guervós. 2010. EvAg: a scalable peer-to-peer evolutionary algorithm. *Genetic Programming and Evolvable Machines* 11, 2 (2010), 227–246.
- [12] Kittisak Kerdprasop and Nittaya Kerdprasop. 2012. Concurrent Computation for Genetic Algorithms. In *Proceedings of the 1st International Conference on Software Technology*. 79–84.
- [13] J.L.J. Laredo, P.A. Castillo, A.M. Mora, and J.J. Merelo. 2008. Exploring Population Structures for Locally Concurrent and Massively Parallel Evolutionary Algorithms. In *WCCI 2008 Proceedings*. IEEE Press, 2610–2617. http://atc.ugr.es/I+D+i/congresos/2008/CEC_2008_2610.pdf
- [14] Juan Luis Jiménez Laredo, Pascal Bouvry, Sanaz Mostaghim, and Juan Julián Merelo Guervós. 2012. Validating a Peer-to-Peer Evolutionary Algorithm. In *Applications of Evolutionary Computation - EvoApplications 2012, Málaga, Spain, April 11-13, 2012, Proceedings*. 436–445. https://doi.org/10.1007/978-3-642-29178-4_44
- [15] Moritz Lenz. 2017. *Perl 6 Fundamentals*. Springer.
- [16] Xia Li, Kunqi Liu, Lixiao Ma, and Huanzhe Li. 2010. A Concurrent-Hybrid Evolutionary Algorithms with Multi-child Differential Evolution and Guotao Algorithm Based on Cultural Algorithm Framework. In *Advances in Computation and Intelligence*, Zhihua Cai, Chengyu Hu, Zhuo Kang, and Yong Liu (Eds.). Springer Berlin Heidelberg, Berlin, Heidelberg, 123–133.

- [16] Juan Julián Merelo Guervós and José Mario García Valdez. 2018. Performance improvements of evolutionary algorithms in perl 6. In *Proceedings of the Genetic and Evolutionary Computation Conference Companion, GECCO 2018, Kyoto, Japan, July 15–19, 2018*, Hernán E. Aguirre and Keiki Takadama (Eds.). ACM, 1371–1378. <https://doi.org/10.1145/3205651.3208273>
- [17] Hans Schippers, Tom Van Cutsem, Stefan Marr, Michael Haupt, and Robert Hirschfeld. 2009. Towards an actor-based concurrent machine model. In *Proceedings of the 4th workshop on the Implementation, Compilation, Optimization of Object-Oriented Languages and Programming Systems*. ACM, 4–9.
- [18] Gene I. Sher. 2013. *Handbook of Neuroevolution Through Erlang*. Springer.
- [19] Jed Simson and Michael Mayo. 2017. Open-Source Linear Genetic Programming. (2017).
- [20] Jerry Swan, Steven Adriaensen, Mohamed Bishr, Edmund K Burke, John A Clark, Patrick De Causmaecker, Juanjo Durillo, Kevin Hammond, Emma Hart, Colin G Johnson, et al. 2015. A research agenda for metaheuristic standardization. In *Proceedings of the XI Metaheuristics International Conference*.
- [21] K. Tagawa. 2012. Concurrent differential evolution based on generational model for multi-core CPUs. *Lecture Notes in Computer Science (including subseries Lecture Notes in Artificial Intelligence and Lecture Notes in Bioinformatics)* 7673 LNCS (2012), 12–21. cited By (since 1996) 0.
- [] Audrey Tang. 2007. Perl 6: Reconciling the Irreconcilable. In *Proceedings of the 34th Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages (POPL '07)*. ACM, New York, NY, USA, 1–1. <https://doi.org/10.1145/1190216.1190218>
- [22] Same A. U. Thors. 2018. Some title. In *Proceedings*. ACM.
- [23] Same A. U. Thors. 2018. Some title. In *Proceedings*. Springer.
- [24] Same A. U. Thors. 2018. Some title. In *Proceedings*. ACM.
- [25] Lazar Valkov, Dipak Chaudhari, Akash Srivastava, Charles Sutton, and Swarat Chaudhuri. 2018. Synthesis of Differentiable Functional Programs for Lifelong Learning. *arXiv preprint arXiv:1804.00218* (2018).
- [26] Paul Walsh. 1999. A Functional Style and Fitness Evaluation Scheme for Inducting High Level Programs. In *Proceedings of the Genetic and Evolutionary Computation Conference*, Wolfgang Banzhaf, Jason Daida, Agoston E. Eiben, Max H. Garzon, Vasant Honavar, Mark Jakiela, and Robert E. Smith (Eds.), Vol. 2. Morgan Kaufmann, Orlando, Florida, USA, 1211–1216. <http://www.cs.bham.ac.uk/~wbl/biblio/gecco1999/GP-455.ps>